

Домашнее задание для к уроку 7 - Продвинутые абстракции Kubernetes

! Задание нужно выполнять в нэймспэйсе default

Разверните в кластере сервер систему мониторинга Prometheus.

- Создайте в кластере ConfigMap со следующим содержимым:

```
prometheus.yml: |
  global:
    scrape_interval: 30s

  scrape_configs:
    - job_name: 'prometheus'
      static_configs:
        - targets: ['localhost:9090']
    - job_name: 'kubernetes-nodes'
      kubernetes_sd_configs:
        - role: node
      relabel_configs:
        - source_labels: [__address__]
          regex: (.+):(.+)
          target_label: __address__
          replacement: ${1}:9101
```

Создайте объекты для авторизации Prometheus сервера в Kubernetes-API.

```
---
apiVersion: v1
kind: ServiceAccount
metadata:
  name: prometheus
  namespace: default
---
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRole
metadata:
  name: prometheus
rules:
- apiGroups: ["" ]
  resources:
    - nodes
  verbs: ["get", "list", "watch"]
---
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRoleBinding
metadata:
  name: prometheus
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: prometheus
subjects:
- kind: ServiceAccount
  name: prometheus
  namespace: default
```

- Создайте StatefulSet для Prometheus сервера из образа prom/prometheus:v2.19.2 с одной репликой

В нем должен быть описан порт 9090 TCP
volumeClaimTemplate - ReadWriteOnce, 5Gi, подключенный по пути /prometheus
Подключение конфигмапа с настройками выше по пути /etc/prometheus

Так же в этом стейтфулсете нужно объявить initContainer для изменения прав на вольюм data на 777.
См пример из лекции 4: practice/4.resources-and-persistence/persistence/deployment.yaml

Не забудьте указать обязательное поле serviceName

Так же укажите поле serviceAccount: prometheus на одном уровне с containers, initContainers, volumes

См пример с rabbitmq из материалов лекции.

- Создайте service и ingress для этого стейтфулсета, так чтобы запросы с любым доменом на белый IP вашего сервиса nginx-ingress-controller (тот что в нэймспэйсе ingress-nginx с типом LoadBalancer) шли на приложение
- Проверьте что при обращении из браузера на белый IP вы видите открывшееся приложение Prometheus
- В этом же неймспэйсе создайте DaemonSet node-exporter как в примере к лекции: practice/7.advanced-abstractions/daemonset.yaml
- Откройте в браузере интерфейс Prometheus.
Попробуйте открыть Status -> Targets
Тут вы должны увидеть все ноды своего кластера, которые Prometheus смог определить и собирает с ним метрики.

Так же можете попробовать на вкладке Graph выполнить запрос node_load1 - это минутный Load Average для каждой из нод в кластере.

ЛИСТИНГ манифестов

```
cat configmap_prometheus.yaml
```

```
---
```

```
apiVersion: v1
```

```
kind: ConfigMap
```

```
metadata:
```

```
  name: prometheus-config
```

```
data:
```

```
  prometheus.yml: |
```

```
    global:
```

```
      scrape_interval: 30s
```

```
  scrape_configs:
```

```
    - job_name: 'prometheus'
```

```
      static_configs:
```

```
        - targets: ['localhost:9090']
```

```
    - job_name: 'kubernetes-nodes'
```

```
      kubernetes_sd_configs:
```

```
        - role: node
```

```
      relabel_configs:
```

```
        - source_labels: [__address__]
```

```
          regex: (.+):(.+)
```

```
          target_label: __address__
```

```
          replacement: ${1}:9101
```

```
cat daemonset.yaml
```

```
---
```

```
apiVersion: apps/v1
```

```
kind: DaemonSet
```

```
metadata:
```

```
  labels:
```

```
    app: node-exporter
```

```
  name: node-exporter
```

```

spec:
  updateStrategy:
    rollingUpdate:
      maxUnavailable: 1
    type: RollingUpdate
  selector:
    matchLabels:
      app: node-exporter
  template:
    metadata:
      labels:
        app: node-exporter
    spec:
      containers:
        - args:
            - --web.listen-address=0.0.0.0:9101
            - --path.procfs=/host/proc
            - --path.sysfs=/host/sys
            - --collector.filesystem.ignored-mount-points=^/(dev|proc|sys|var/lib/
docker/.+)($/|/)
            - --collector.filesystem.ignored-fs-types=^(autofs|binfmt_misc|cgroup|
configfs|debugfs|devpts|devtmpfs|fusectl|hugetlbfs|mqueue|overlay|proc|procfs|
pstore|rpc_pipefs|securityfs|sysfs|tracefs)$
          image: quay.io/prometheus/node-exporter:v0.16.0
          imagePullPolicy: IfNotPresent
          name: node-exporter
          volumeMounts:
            - mountPath: /host/proc
              name: proc
            - mountPath: /host/sys
              name: sys
            - mountPath: /host/root
              name: root
              readOnly: true
          hostNetwork: true

```

```
hostPID: true
tolerations:
  - effect: NoSchedule
    operator: Exists
nodeSelector:
  beta.kubernetes.io/os: linux
volumes:
  - hostPath:
      path: /proc
      type: ""
    name: proc
  - hostPath:
      path: /sys
      type: ""
    name: sys
  - hostPath:
      path: /
      type: ""
    name: root
```

cat ingress_prometheus.yaml

```
apiVersion: extensions/v1beta1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: prometheus
```

```
spec:
```

```
  rules:
```

```
    - http:
```

```
paths:
  - path: /
    pathType: Prefix
    backend:
      serviceName: prometheus
      servicePort: 80
```

```
cat service_acc_prometheus.yaml
```

```
---
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: prometheus
  namespace: default
```

```
---
```

```
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRole
metadata:
  name: prometheus
rules:
  - apiGroups: [""]
    resources:
```

```
- nodes

verbs: ["get", "list", "watch"]

---

apiVersion: rbac.authorization.k8s.io/v1beta1

kind: ClusterRoleBinding

metadata:

  name: prometheus

roleRef:

  apiGroup: rbac.authorization.k8s.io

  kind: ClusterRole

  name: prometheus

subjects:

- kind: ServiceAccount

  name: prometheus

  namespace: default
```

cat service_prometheus.yaml

```
---

apiVersion: v1

kind: Service

metadata:

  name: prometheus
```

```
spec:

  selector:

    app: prometheus

  ports:

  - protocol: TCP

    port: 80

    targetPort: 9090
```

```
cat sts_prometheus3.yaml
```

```
---
```

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: prometheus
spec:
  serviceName: "prometheus"
  replicas: 1
  selector:
    matchLabels:
      app: prometheus
  template:
    metadata:
      labels:
        app: prometheus
    spec:
```



```
serviceAccount: prometheus

initContainers:
- name: "init-busybox-permission777"
  image: busybox
  imagePullPolicy: Always
  command: ["chmod", "-R", "777", "/data"]
  volumeMounts:
  - name: prometheus-data
    mountPath: /data

containers:
- name: prometheus-server-2-19-2
  image: prom/prometheus:v2.19.2
  imagePullPolicy: "IfNotPresent"
  ports:
  - containerPort: 9090
  readinessProbe:
    httpGet:
      path: /-/ready
      port: 9090
    initialDelaySeconds: 30
    timeoutSeconds: 30
  livenessProbe:
    httpGet:
      path: /-/healthy
      port: 9090
    initialDelaySeconds: 30
    timeoutSeconds: 30
```

resources:

requests:

cpu: 200m

memory: 1000Mi

volumeMounts:

- name: config-volume-prometheus

mountPath: /etc/prometheus

- name: prometheus-data

mountPath: /data

terminationGracePeriodSeconds: 300

volumes:

- name: config-volume-prometheus

configMap:

name: prometheus-config

volumeClaimTemplates:

- metadata:

name: prometheus-data

spec:

storageClassName: csi-ceph-hdd-dp1

accessModes:

- ReadWriteOnce

resources:

requests:

storage: "5Gi"

Скриншоты работы kubernetes-cluster

```
lamo@lamo-K501LB:~/nout/GEEKBRAINS/Microservice_architecture_and_containerization/HW7$ k get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
prometheus-data-prometheus-0	Bound	pvc-7b64bf0d-15e7-4872-afef-031668abb7c0	5Gi	RWO	csi-ceph-hdd-dp1	17s
ps-pgdata	Bound	pvc-183ae07a-2319-4ea5-a132-91fa99a8e6ab	10Gi	RWX	csi-ceph-hdd-dp1	6d17h
webdav	Bound	pvc-40024748-9194-4a9b-a5f9-fe7837df9870	2Gi	RWX	csi-ceph-hdd-dp1	9d

```
lamo@lamo-K501LB:~/nout/GEEKBRAINS/Microservice_architecture_and_containerization/HW7$ k get svc -A
```

NAMESPACE	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	kubernetes	ClusterIP	10.254.0.1	<none>	443/TCP	10d
default	prometheus	ClusterIP	10.254.141.68	<none>	80/TCP	37s
ingress-nginx	nginx-ingress-controller	LoadBalancer	10.254.128.123	89.208.87.242	80:30080/TCP,443:30443/TCP	10d
ingress-nginx	nginx-ingress-controller-metrics	ClusterIP	10.254.11.207	<none>	9913/TCP	10d
ingress-nginx	nginx-ingress-default-backend	ClusterIP	10.254.244.136	<none>	80/TCP	10d
kube-system	calico-node	ClusterIP	None	<none>	9091/TCP	10d
kube-system	csi-cinder-controller-service	ClusterIP	10.254.209.85	<none>	12345/TCP	10d
kube-system	dashboard-metrics-scraper	ClusterIP	10.254.174.238	<none>	8000/TCP	10d
kube-system	kube-dns	ClusterIP	10.254.0.10	<none>	53/UDP,53/TCP,9153/TCP	10d
kube-system	kubernetes-dashboard	ClusterIP	10.254.179.232	<none>	443/TCP	10d
kube-system	metrics-server	ClusterIP	10.254.209.248	<none>	443/TCP	10d
magnum-tiller	tiller-deploy	ClusterIP	10.254.126.244	<none>	44134/TCP	10d

```
lamo@lamo-K501LB:~/nout/GEEKBRAINS/Microservice_architecture_and_containerization/HW7$ k get sts
```

NAME	READY	AGE
prometheus	1/1	67s

```
lamo@lamo-K501LB:~/nout/GEEKBRAINS/Microservice_architecture_and_containerization/HW7$ k get ing
```

NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
prometheus	<none>	*		80	72s

```
lamo@lamo-K501LB:~/nout/GEEKBRAINS/Microservice_architecture_and_containerization/HW7$ k get ep
```

NAME	ENDPOINTS	AGE
kubernetes	10.0.0.9:6443	10d
prometheus	10.100.180.144:9090	87s

```
lamo@lamo-K501LB:~/nout/GEEKBRAINS/Microservice_architecture_and_containerization/HW7$ k get po -owide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
node-exporter-5295h	1/1	Running	0	113s	10.0.0.9	kubernetes-cluster-5268-master-0	<none>	<none>
node-exporter-cdklx	1/1	Running	0	113s	10.0.0.14	kubernetes-cluster-5268-default-group-0	<none>	<none>
node-exporter-dpm56	1/1	Running	0	113s	10.0.0.19	kubernetes-cluster-5268-default-group-2	<none>	<none>
prometheus-0	1/1	Running	0	113s	10.100.180.144	kubernetes-cluster-5268-default-group-0	<none>	<none>

Подключаемся через браузер к нашему приложению.

https://89.208.87.242/targets

☆

Prometheus

Alerts

Graph

Status ▾

Help

Targets

All

Unhealthy

kubernetes-nodes (3/3 up)

show less

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://10.0.0.14:9101/metrics	UP	<div>instance="kubernetes-cluster-5268-default-group-0"</div> <div>job="kubernetes-nodes"</div>	5.407s ago	25.04ms	
http://10.0.0.19:9101/metrics	UP	<div>instance="kubernetes-cluster-5268-default-group-2"</div> <div>job="kubernetes-nodes"</div>	20.071s ago	22.81ms	
http://10.0.0.9:9101/metrics	UP	<div>instance="kubernetes-cluster-5268-master-0"</div> <div>job="kubernetes-nodes"</div>	13.764s ago	22.47ms	

prometheus (1/1 up)

show less

Endpoint	State	Labels	Last Scrape	Scrape Duration	Error
http://localhost:9090/metrics	UP	<div>instance="localhost:9090"</div> <div>job="prometheus"</div>	4.183s ago	9.024ms	

Prometheus

Alerts

Graph

Status ▾

Help

Enable query history

node_load1

Execute

- insert metric at cursor - ▾

Graph

Console

◀

Moment

▶

Element	Value
node_load1{instance="kubernetes-cluster-5268-master-0",job="kubernetes-nodes"}	0.37
node_load1{instance="kubernetes-cluster-5268-default-group-0",job="kubernetes-nodes"}	0.02
node_load1{instance="kubernetes-cluster-5268-default-group-2",job="kubernetes-nodes"}	0.37

Try experimental React UI

Load time: 90ms

Resolution: 14s

Total time series: 3